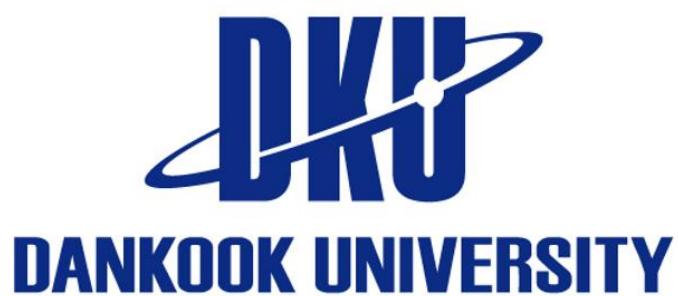


RocksDB on Ubuntu

Quick Start Guide

Dankook Univ. System Software Lab

Dayeon Wee, Yongmin Lee

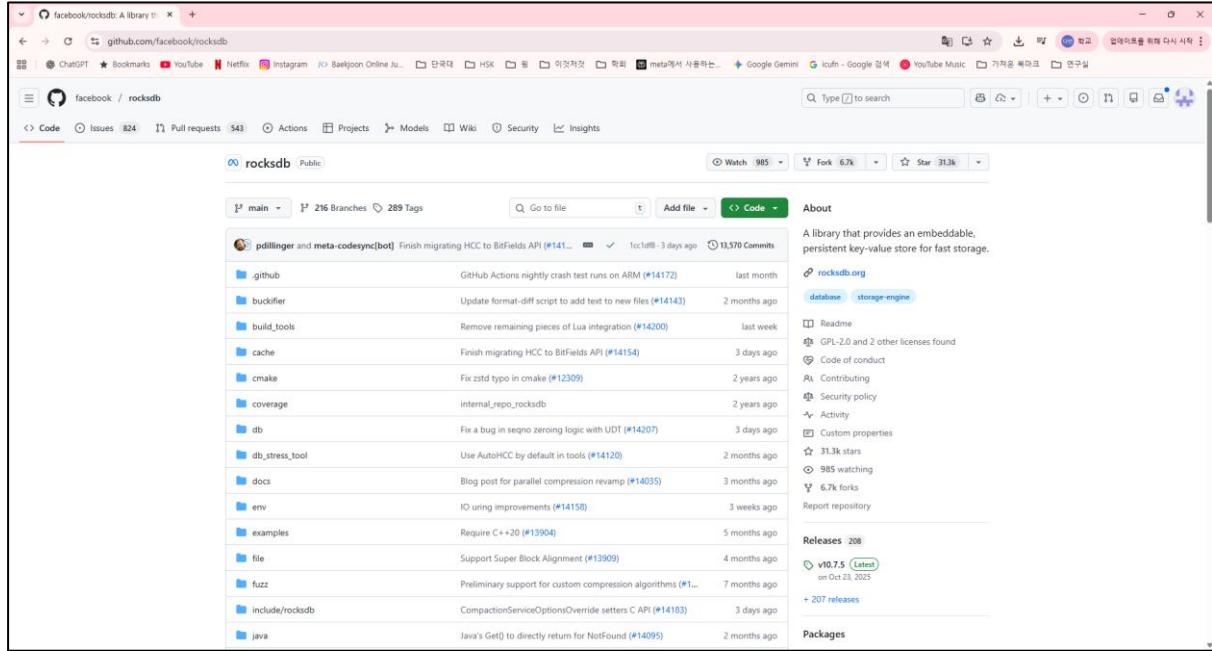


목차

1. RocksDB 설치	3
2. RocksDB 빌드	4
2.1 의존성 설치.....	4
2.2 빌드.....	4
3. db_bench 실행.....	4
3.1 db_bench 옵션	5
4. 진행 방법	9
4.1 코드 분석 및 수정	9
4.2 옵션 튜닝	9

1. RocksDB 설치

<https://github.com/facebook/rocksdb>



우분투 환경에서

git clone <https://github.com/facebook/rocksdb>

```
● (base) dy@choi-gunhee-linux-93:~/workspace$ git clone https://github.com/facebook/rocksdb
Cloning into 'rocksdb'...
remote: Enumerating objects: 142596, done.
remote: Counting objects: 100% (130/130), done.
remote: Compressing objects: 100% (91/91), done.
remote: Total 142596 (delta 80), reused 39 (delta 39), pack-reused 142466 (from 3)
Receiving objects: 100% (142596/142596), 227.52 MiB | 20.12 MiB/s, done.
Resolving deltas: 100% (109457/109457), done.
```

cd RocksDB

ls

RocksDB 내부 파일들이 위치해 있음

```
● (base) dy@choi-gunhee-linux-93:~/workspace$ cd ./rocksdb/
● (base) dy@choi-gunhee-linux-93:~/workspace/rocksdb$ ls
AUTHORS           common.mk          docs              INSTALL.md      memtable        src.mk          util
BUCK             CONTRIBUTING.md    DUMP_FORMAT.md   issue_template.md  microbench     table
buckifier        COPYING           env               java            monitoring    test_util      utilities
build_tools      coverage          examples         LANGUAGE-BINDINGS.md options        third-party  Vagrantfile
cache            crash_test.mk     file              LICENSE.Apache plugin         tools
ccache_msvc_compiler.bat db               fuzz             LICENSE.leveldb PLUGINS.md   thirdparty.inc WINDOWS_PORT.md
cmake            db_stress_tool    logging          Makefile        README.md    trace_replay
CMakeLists.txt    DEFAULT_OPTIONS_HISTORY.md memory          rocksdb.pc.in  unreleased_history
CODE_OF_CONDUCT.md Directory.Build.props include          USERS.md
```

2. RocksDB 빌드

2.1 의존성 설치

```
sudo apt-get install libgflags-dev
```

```
sudo apt-get install libsnappy-dev
```

```
sudo apt-get install zlib1g-dev
```

```
sudo apt-get install libbz2-dev
```

```
sudo apt-get install liblzl4-dev
```

```
sudo apt-get install libzstd-dev
```

2.2 빌드

```
make db_bench
```

```
(base) dy@choi-gunhee-linux-93:~/workspace/rocksdb$ make db_bench
$DEBUG_LEVEL is 1, $LIB_MODE is shared
Makefile:184: Warning: Compiling in debug mode. Don't use the resulting binary in production
$DEBUG_LEVEL is 1, $LIB_MODE is shared
Makefile:184: Warning: Compiling in debug mode. Don't use the resulting binary in production
CC      tools/db_bench.o
CC      tools/db_bench_tool.o
CC      tools/tool_hooks.o
CC      tools/simulated_hybrid_file_system.o
CC      test_util/testutil.o
CC      cache/cache.o
CC      cache/cache_entry_roles.o
CC      cache/cache_key.o
CC      cache/cache_helpers.o
CC      cache/cache_reservation_manager.o
CC      cache/charged_cache.o
CC      cache/clock_cache.o
CC      cache/lru_cache.o
CC      cache/compressed_secondary_cache.o
CC      cache/secondary_cache.o
CC      cache/secondary_cache_adapter.o
CC      cache/sharded_cache.o
```

빌드 시 약 5~10분 소요

3. db_bench 실행

- db_bench는 RocksDB에 포함된 성능 벤치마크 및 스트레스 테스트 도구
- 다양한 워크로드를 실행해 지연시간·처리량·쓰기 증폭을 측정

- RocksDB 설정 옵션을 바꿔가며 튜닝 효과를 정량적으로 비교 가능

./db_bench를 모든 옵션을 디폴트로 실행했을 때의 출력 결과

```
④ (base) dy@choi-gunhee-linux-93:~/workspace/rocksdb$ ./db_bench
Set seed to 1767340342791125 because --seed was 0
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
Integrated BlobDB: blob cache disabled
RocksDB:      version 10.11.0
Date:        Fri Jan  2 16:52:22 2026
CPU:         8 * Intel(R) Core(TM) i7-6700 CPU @ 3.40GHz
CPUCache:    8192 KB
Keys:        16 bytes each (+ 0 bytes user-defined timestamp)
Values:      100 bytes each (50 bytes after compression)
Entries:     1000000
Prefix:      0 bytes
Keys per prefix:   0
RawSize:    110.6 MB (estimated)
FileSize:   62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: SkipListFactory
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
Integrated BlobDB: blob cache disabled
DB path: [/tmp/rocksdbtest-1001/dbbench]
fillseq : 3.448 micros/op 290040 ops/sec 3.448 seconds 1000000 operations; 32.1 MB/s
Please disable_auto_compactions in FillDeterministic benchmark
```

3.1 db_bench 옵션

기본적인 옵션들은 위키에 있음

<https://github.com/facebook/rocksdb/wiki/Benchmarking-tools>

The screenshot shows a GitHub repository page for 'rocksdb' under the 'facebook' organization. The page title is 'Benchmarking tools'. The main content is about the 'db_bench' tool, which is described as the main tool for benchmarking RocksDB's performance. It supports many benchmarks to generate different types of workloads. The page includes examples of command-line usage, such as './db_bench --benchmarks="fillseq" --statistics' and './db_bench --benchmarks="fillseq,stats" --statistics'. A sidebar on the right lists various 'Contents' related to RocksDB, including 'RocksDB Wiki', 'Overview', 'Contributor FAQ', 'Terminology', 'Requirements', 'Release Methodology', 'RocksDB Users and Use Cases', 'RocksDB Public Communication and Information Channels', and several sub-sections like 'Iterators', 'Prefix seek', 'SeekForPrev', 'Tailing Iterator', 'Compaction Filter', 'Multi Column Family Iterator', 'Read-Modify-Write (Merge) Operator', 'Column Family', 'Creating and Ingesting SST File', 'Single Delete', 'SST Partitioner', 'Low Priority Write', 'Time to Live (TTL) Support', 'Transactions', 'Snapshot', 'DeleteRange', and 'Atomicity'.

예를 들어,

`./db_bench --benchmarks="fillseq,stats" --statistics` 와 같이 실행하게 되면

```
(base) dy@choi-gunhee-linux-93:~/workspace/rocksdb$ ./db_bench --benchmarks="fillseq,stats" --statistics
Set seed to 1767341446843694 because --seed was 0
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
Integrated BlobDB: blob cache disabled
RocksDB:    version 10.11.0
Date:      Fri Jan 2 17:10:46 2026
CPU:       8 * Intel(R) Core(TM) i7-6700 CPU @ 3.40GHz
CPUCache:   8192 KB
Keys:      16 bytes each (+ 0 bytes user-defined timestamp)
Values:    100 bytes each (50 bytes after compression)
Entries:   1000000
Prefix:    0 bytes
Keys per prefix: 0
RawSize:   110.6 MB (estimated)
FileSize:  62.9 MB (estimated)
Write rate: 0 bytes/second
Read rate: 0 ops/second
Compression: Snappy
Compression sampling rate: 0
Memtablerep: SkipListFactory
Perf Level: 1
WARNING: Assertions are enabled; benchmarks unnecessarily slow
-----
Initializing RocksDB Options from the specified file
Initializing RocksDB Options from command-line flags
Integrated BlobDB: blob cache disabled
DB path: [/tmp/rocksdbtest-1001/dbbench]
fillseq     :      3.751 micros/op 266573 ops/sec 3.751 seconds 1000000 operations;  29.5 MB/s
```

간략한 실행 결과와

```

** Compaction Stats [default] **
Level  Files  Size   Score Read(GB) Rm(GB) Rmp(GB) Write(GB) WPreComp(GB) Wnew(GB) Moved(GB) W-Amp Rd(MB/s) Comp(sec) CompMergeCPU(sec) Comp(cnt) Avg(sec) KeyIn KeyDrop Rblob(GB) Wblob(GB)
L0    2/0   59.28 MB 0.5   0.0   0.0   0.0   0.1   0.1   0.1   0.0   1.0   0.0   108.7   0.55   0.39   2   0.273   921K   0   0.0   0.0
Sum   2/0   59.28 MB 0.0   0.0   0.0   0.0   0.1   0.1   0.1   0.0   1.0   0.0   108.7   0.55   0.39   2   0.273   921K   0   0.0   0.0
Int   0/0   0.00 KB 0.0   0.0   0.0   0.0   0.1   0.1   0.1   0.0   1.0   0.0   108.7   0.55   0.39   2   0.273   921K   0   0.0   0.0

** Compaction Stats [default] **
Priority  Files  Size   Score Read(GB) Rm(GB) Rmp(GB) Write(GB) WPreComp(GB) Wnew(GB) Moved(GB) W-Amp Rd(MB/s) Comp(sec) CompMergeCPU(sec) Comp(cnt) Avg(sec) KeyIn KeyDrop Rblob(GB) Wblob(GB)
High     0/0   0.00 KB 0.0   0.0   0.0   0.0   0.1   0.1   0.1   0.0   0.0   0.0   108.7   0.55   0.39   2   0.273   921K   0   0.0   0.0

Blk file count: 0, total size: 0.0 GB, garbage size: 0.0 GB, space amp: 0.0

Uptime(sec): 3.8 total, 3.8 interval
Flush(GB): cumulative 0.058, interval 0.058
AddFile(GB): cumulative 0.008, interval 0.008
AddFile(1MB): cumulative 0.008, interval 0
AddFile(8MB): cumulative 0.008, interval 0
AddFile(64MB): cumulative 0.008, interval 0
AddFile(Keys): cumulative 0, interval 0

Cumulative compaction: 0.00 GB write, 15.74 MB/s write, 0.00 GB read, 0.00 MB/s read, 0.5 seconds
Interval compaction: 0.00 GB write, 15.74 MB/s write, 0.00 GB read, 0.00 MB/s read, 0.5 seconds
Estimated bytes written by compaction: 0
Write Stall (count): cf-10-file-count-limit-delays: 0, cf-10-file-count-limit-stops-with-ongoing-compaction: 0, cf-10-file-count-limit-stops: 0, memtable-limit-delays: 0, memtable-limit-stops: 0, pending-compaction-bytes-delays: 0, pending-compaction-bytes-stops: 0, total-delays: 0, total-stops: 0
Block cache AutoclockCache@0xfef5a24769092233400 Capacity: 32.00 MB seed: 10595245358 usage: 4.00 KB table_size: 64 occupancy: 1 collections: 1 last_copies: 0 last_secs: 2.1e-05 secs_since: 3
Block cache entry stats(count,size,portion): Misc(1,0.00 KB,0%)

** File Read Latency Histogram [default] **
** Level 0 read latency histogram (micros):**
Count: 1 Average: 8.87 Median: 8.87 Max: 23
Percentiles: P50: 2.00 P75: 10.00 P90: 23.00 P99.0: 23.00 P99.9: 23.00

[  0,     1 ]      2  25.0000 #####*
[  1,     2 ]      2  25.0000 #####*
(  4,     6 ]      1  12.5000 62.5000 ####
(  6,    10 ]      1  12.5000 75.0000 ####
( 22,    34 ]      2  25.0000 100.0000 ####

** DB Stats **
Uptime(sec): 3.8 total, 3.8 interval
Cumulative writes: 1000K writes, 1000K keys, 1000K commit groups, 1.0 writes per commit group, ingest: 0.12 GB, 33.18 MB/s
Cumulative WAL: 1000K writes, 0 syncs, 1000000.00 writes per sync, written: 0.12 GB, 33.18 MB/s
Cumulative stall: 00:00:0.000 H:M:S, 0.0 percent
Interval writes: 1000K writes, 1000K keys, 1000K commit groups, 1.0 writes per commit group, ingest: 124.93 MB, 33.18 MB/s
Interval WAL: 1000K writes, 0 syncs, 1000000.00 writes per sync, written: 0.12 GB, 33.18 MB/s
Interval stall: 00:00:0.000 H:M:S, 0.0 percent
Write Stall (count): write-buffer-manager-limit-stops: 0

STATISTICS:
rocksdb.block.cache.miss COUNT : 0
rocksdb.block.cache.hit COUNT : 0
rocksdb.block.cache.add COUNT : 0
rocksdb.block.cache.add.failures COUNT : 0
rocksdb.block.cache.index.miss COUNT : 0
rocksdb.block.cache.index.hit COUNT : 0
rocksdb.block.cache.index.add COUNT : 0
rocksdb.block.cache.index.bytes.insert COUNT : 0
rocksdb.block.cache.filter.miss COUNT : 0
rocksdb.block.cache.filter.hit COUNT : 0
rocksdb.block.cache.filter.add COUNT : 0
rocksdb.block.cache.filter.bytes.insert COUNT : 0
rocksdb.block.cache.data.miss COUNT : 0
rocksdb.block.cache.data.hit COUNT : 0
rocksdb.block.cache.data.add COUNT : 0
rocksdb.block.cache.data.bytes.insert COUNT : 0
rocksdb.block.cache.bytes.read COUNT : 0
rocksdb.block.cache.bytes.write COUNT : 0
rocksdb.block.cache.compression.dict.miss COUNT : 0
rocksdb.block.cache.compression.dict.hit COUNT : 0
rocksdb.block.cache.compression.dict.add COUNT : 0
rocksdb.block.cache.compression.dict.bytes.insert COUNT : 0
rocksdb.block.cache.add.redundant COUNT : 0
rocksdb.block.cache.index.add.redundant COUNT : 0
rocksdb.block.cache.filter.add.redundant COUNT : 0
rocksdb.block.cache.data.add.redundant COUNT : 0
rocksdb.block.cache.compression.dict.add.redundant COUNT : 0
rocksdb.secondary.cache.hits COUNT : 0
rocksdb.secondary.cache.filter.hits COUNT : 0
rocksdb.secondary.cache.index.hits COUNT : 0
rocksdb.secondary.cache.data.hits COUNT : 0
rocksdb.compressed.secondary.cache.dummy.hits COUNT : 0
rocksdb.compressed.secondary.cache.hits COUNT : 0
rocksdb.compressed.secondary.cache.promotions COUNT : 0
rocksdb.compressed.secondary.cache.promotion.skips COUNT : 0

```

Benchmarks=stats와 statistics 옵션을 줌으로써 더 자세한 결과를 확인할 수 있음

또한 터미널에서 ./db_bench --help 를 입력하면 db_bench 에서 사용할 수 있는 여러 다양한 옵션들을 출력해줌

따로 옵션들에 대해 설정하지 않은 값은 default: xx 로 지정되어 있어서 자동으로 default 값 사용

```
(base) dy@choi-gunhee-linux-93:~/workspace/rocksdb$ ./db_bench --help
  -data_block_hash_table_util_ratio (util ratio for data block hash index
    table. This is only valid if use_data_block_hash_index is set to true)
    type: double default: 0.75
  -db (Use the db with the following name.) type: string default: ""
  -db_write_buffer_size (Number of bytes to buffer in all memtables before
    compacting) type: int64 default: 0
  -decouple_partitioned_filters (Decouple filter partitioning from index
    partitioning.) type: bool default: true
  -delayed_write_rate (Limited bytes allowed to DB when soft_rate_limit or
    level0_slowdown_writes_trigger triggers) type: uint64 default: 8388608
  -delete_obsolete_files_period_micros (Ignored. Left here for backward
    compatibility) type: uint64 default: 0
  -deletepercent (Percentage of deletes out of reads/writes/deletes (used in
    RandomWithVerify only). RandomWithVerify calculates writepercent as (100
    - FLAGS_readwritepercent - deletepercent), so deletepercent must be
    smaller than (100 - FLAGS_readwritepercent)) type: int32 default: 2
  -deletes (Number of delete operations to do. If negative, do FLAGS_num
    deletions.) type: int64 default: -1
  -disable_auto_compactions (Do not auto trigger compactions) type: bool
    default: false
  -disable_seek_compaction (Not used, left here for backwards compatibility)
    type: int32 default: 0
  -disable_wal (If true, do not write WAL for write.) type: bool
    default: false
  -disposable_entries_batch_size (Number of consecutively inserted disposable
    KV entries that will be deleted after 'delete_delay' microseconds. A
    series of Deletes is always issued once all the disposable KV entries it
    targets have been inserted into the DB. When 0 no deletes are issued and
    a regular 'fillanddeleteuniquerandom' benchmark occurs. (only compatible with
    fillanddeleteuniquerandom benchmark)) type: uint64 default: 0
  -disposable_entries_delete_delay (Minimum delay in microseconds for the
    series of Deletes to be issued. When 0 the insertion of the last
    disposable entry is immediately followed by the issuance of the Deletes.
    (only compatible with fillanddeleteuniquerandom benchmark).) type: uint64
    default: 0
  -disposable_entries_value_size (Size of the values (in bytes) of the
    entries targeted by selective deletes. (only compatible with
    fillanddeleteuniquerandom benchmark)) type: int32 default: 64
  -dump_malloc_stats (Dump malloc stats in LOG ) type: bool default: true
  -duration (Time in seconds for the random-ops tests to run. When 0 then num
    & reads determine the test duration) type: int32 default: 0
  -enable_blob_files ([Integrated BlobDB] Enable writing large values to
    separate blob files.) type: bool default: false
  -enable_blob_garbage_collection ([Integrated BlobDB] Enable blob garbage
    collection.) type: bool default: false
  -enable_cpu_prio (Lower the background flush/compaction threads' CPU
    priority) type: bool default: false
  -enable_index_compression (Compress the index block) type: bool
    default: true
  -enable_io_prio (Lower the background flush/compaction threads' IO
    priority) type: bool default: false
```

캡쳐된 부분 이외에도 100개가 넘는 옵션들이 존재함

4. 진행 방법

4.1 코드 분석 및 수정

https://github.com/DKU-StarLab/2026_RocksDB_Study/tree/main/presentation_file

깃허브에 업로드된 W0_DayeonWee _How_To_Analyze_RocksDB 자료를 참고

또한 lm에게 물어보는 것도 방법이긴 하나, Hallucination이 있으니 너무 믿진 말 것

4.2 옵션 튜닝

사진에 나와있는 것처럼 옵션을 변경할 수 있음

Experiment 1 - Measurement setup (modified)

1. 데이터 크기를 늘림

- 기존: `--num=1M` (100만개, default), `--value_size=100` (100B, default)
- 변경: `--num=10M, --value_size=4096`

2. 1번 적용 후 `compaction_style=3`에서 `db_bench`가 도중에 멈춰버리는 문제 발생

Compaction이 없기에 L0의 파일 개수가 RocksDB의 한계를 초과하여
진행이 불가했던 것으로 추정

- 기존: `compaction_style 0,1,2,3`
- 변경: `compaction_style 0,1,2`

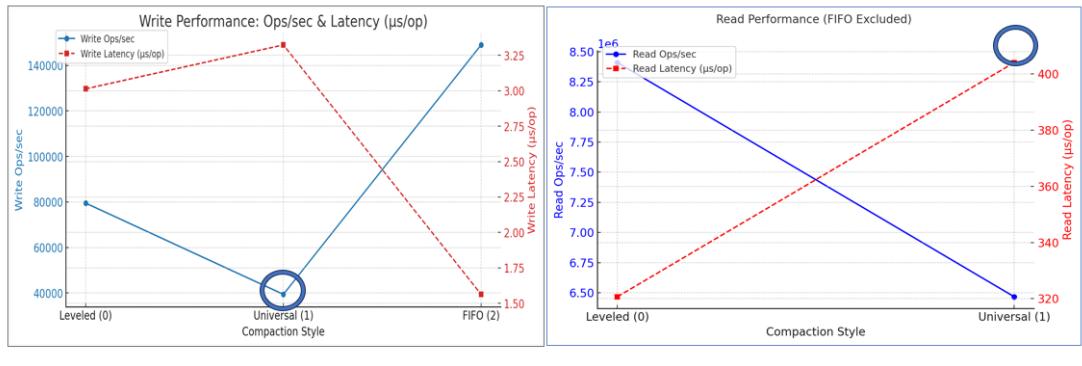
출처: https://github.com/DKU-StarLab/1DanRock/blob/main/presentation_file/Team2/Team2_Week5.pptx

실험을 진행하기 전, 가설을 필수적으로 세워야 함

이 실험을 진행하면 어떠한 이유 때문에 이렇게 결과가 나올 것이다! 와 같은
가설을 세워야 함

또한 이 실험을 왜 진행했는지 motivation이 있으면 좋음

Experiment 1 – Result (graph)



다음과 같이 실험 결과를 그래프로 뽑아볼 수 있음

그래프는 jupyter notebook (python)이나 엑셀을 통해서 그래프를 그려볼 수 있음

이를 통해 결과를 분석해볼 수 있음

화이팅..!